



Texas A&M University - Commerce
Department of Computer Science

Performance Analysis of LR and SVM - Supervised Machine Learning Algorithms for Diabetes Prediction

Vyshnavi Sanikommu

Supervisor: Derek Harter, Ph.D.

A report submitted in partial fulfilment of the requirements of
Texas A&M University - Commerce for the degree of
Master of Science in *Computer Science*

April 24, 2024

Declaration

I, Vyshnavi Sanikommu, of the Department of Computer Science, Texas A&M University - Commerce, confirm that this is my own work and figures, tables, equations, code snippets, artworks, and illustrations in this report are original and have not been taken from any other person's work, except where the works of others have been explicitly acknowledged, quoted, and referenced. I understand that if failing to do so will be considered a case of plagiarism. Plagiarism is a form of academic misconduct and will be penalised accordingly.

I give consent to a copy of my report being shared with future students as an exemplar.

I give consent for my work to be made available more widely to members of TAMUC and public with interest in teaching, learning and research.

Vyshnavi Sanikommu
April 24, 2024

Abstract

Diabetes is one of the most lethal diseases affecting 537 million people worldwide, is projected to rise to 783 million by 2045. Diabetes is a disease caused due to an increase in blood glucose level, causing symptoms like frequent urination, increased hunger, and thirst. Diabetes is a leading cause of blindness, kidney failure, amputations, heart failure and stroke. The body's conversion of food into glucose requires insulin, released by the pancreas, which unlocks cells for glucose entry. This process allows cells to use glucose as an energy source, supporting vital bodily functions. The aim of this project is to develop a system which can perform early prediction of diabetes for a patient, employing supervised machine learning algorithms such as logistic regression and support vector machine. Machine learning techniques provide better results for prediction by constructing models from datasets collected from patients. The project will focus on optimizing performance metrics through a systematic search and tuning of model hyper parameters to maximize overall model effectiveness. The evaluation of model performance will encompass accuracy, precision, recall, f1-score, and confusion matrix. We will compare the performance metrics of the base model with those obtained after tuning using grid search. The results indicate that grid search provides optimal hyper parameters, contributing to the determination of the best-performing model.

Keywords: diabetes, machine learning, logistic regression, support vector machine, performance metrics

Acknowledgements

An acknowledgements section is optional. You may like to acknowledge the support and help of your supervisor(s), friends, or any other person(s), department(s), institute(s), etc. If you have been provided specific facility from department/school acknowledged so.

Contents

1	Introduction	1
1.1	Background	1
1.2	Research question	2
1.3	Aims and objectives	2
1.4	Solution approach	3
1.4.1	Dataset Description	3
1.4.2	Data Preprocessing	3
1.4.3	Model Implementation	3
1.4.4	Hyper Parameter Tuning and Performance Analysis	4
1.5	Summary of contributions and achievements	4
2	Literature Review	5
2.1	Review of state-of-the-art	5
2.2	Critique of the review	6
2.3	Summary	6
3	Methodology	7
3.1	Dataset Description and Data Exploration	7
3.1.1	Data Distribution using Histograms	8
3.1.2	Correlation Matrix	8
3.1.3	Bar plot for Outcome class	9
3.2	Data Preprocessing	10
3.2.1	Missing Values Identification	10
3.2.2	Feature Selection based on Correlation Coefficient	11
3.2.3	Data Normalization	11
3.3	Dataset Split into Train and Test Data	12
3.4	Model Implementation	12
3.4.1	Logistic Regression	12
3.4.2	Support Vector Machine	12
3.4.3	Hyperparameter Tuning with GridSearchCV	13
3.5	Evaluation Metrics	14
3.5.1	Classification Accuracy	14
3.5.2	Confusion Matrix	14
3.5.3	Precision	14
3.5.4	Recall	15

3.5.5 F1-score	15
3.6 Summary	15
4 Results	16
4.1 Results for ML methods - LR and SVM	16
4.2 Hyperparameter Tuning Results for LR and SVM Using GridSearchCV	17
4.3 Summary	17
5 Discussion and Analysis	18
5.1 Discussion on performance metrics	18
5.2 Significance of the findings	18
5.3 Limitations	18
5.4 Summary	19
6 Conclusions and Future Work	20
6.1 Conclusions	20
6.2 Future work	20
7 Reflection	22
Appendices	24
A An Appendix Chapter (Optional)	24
B An Appendix Chapter (Optional)	25

List of Figures

3.1	Top 5 patients data	7
3.2	Data Distribution for each attribute	8
3.3	Correlation Matrix	9
3.4	Distribution of Outcome (0s and 1s)	10
3.5	Top 5 patients data after handling missing values	11

List of Tables

3.1	Dataset attributes and their data types	7
3.2	The number of zero missing values in dataset	10
3.3	The correlation coefficient values	11
4.1	Performance metrics	16
4.2	Confusion Matrix for SVM (Test Data)	16
4.3	Performance metrics after hyperparameter tuning	17
4.4	Confusion Matrix for SVM (Test Data)	17

List of Abbreviations

SMPCS School of Mathematical, Physical and Computational Sciences

Chapter 1

Introduction

Diabetes is a fast growing disease among the people even in youngsters nowadays. It is necessary to understand how it develops in our body. Firstly, we need to understand how a body works without diabetes. Sugar comes from the food that we eat, specially carbohydrates. When we eat this food, the body breaks them down to Sugars or Glucose. This glucose moves around the body in the bloodstream. This glucose is needed by body parts like the brain and pancreas to function. The remainder of glucose is taken to the cells of our body and liver which is stored as energy for later use. In order to use glucose for our body, insulin is required which is a hormone generated by pancreas. Insulin is a key to a closed door which helps glucose moves from blood stream. If pancreas is not able to produce enough insulin or if our body cannot use insulin it produces then glucose levels increases in bloodstream which leads to diabetes.

According to World Health Organization, diabetes is major cause of death worldwide. Around 422 million people worldwide have diabetes. Indeed, it caused deaths of 2 million people in 2019.

There are two types of diabetes present as a disease in human beings: **Type 1** diabetes appear most often during childhood and is characterized by the partial functioning of the pancreas. The cells fail to produce sufficient amounts of Insulin. Initially, we do not see any symptoms as the pancreas remains partially functional. There is no proven study and known methods for prevention. **Type 2** diabetes affects how the body uses sugars for energy. The cells produce low quantity of insulin or the body stops using insulin which can lead to high levels of sugar in bloodstream. High levels of glucose in the bloodstream and urine referred as diabetes mellitus. It is most common type of diabetes found in many people. It is caused by genetic factors and the lifestyle. It affects older adults and more obese or overweight people.

Early prediciton of diabetes can help in controlling the disease and potentially save lives. To accomplish this, this work explores prediction of diabetes by taking various attributes related to diabetes disease. For this purpose, we use the diabetes dataset (Johndasilva, 2018) which has 2000 instances with 9 attributes - 'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome' for creating classification models using logistic regression and support vector machine algorithms.

1.1 Background

This study is undertaken to address the crucial need for accurate and reliable methods in predicting diabetes using logistic regression and support vector machine algorithms on patient data. The

motivation behind this study stems from the growing significance of leveraging machine learning algorithms to assist medical professionals in diagnosing and managing diabetes. Predictive models based on patient data offer the potential to identify individuals at risk or in the early stages of diabetes, enabling proactive and personalized healthcare strategies. Algorithms - logistic regression and support vector machine are chosen due to their widespread use in medical data analysis and their capability to handle classification tasks. The study outcomes have practical implications for healthcare practitioners and researchers, offering insights into algorithm selection for accurate and interpret diabetes prediction.

1.2 Research question

The research questions are:

1. How do the performance metrics of Logistic Regression (LR) and Support Vector Machines (SVM) models differ in predicting diabetes based on patient's data?
2. What factors contribute to these differences, particularly in relation to the hyper parameters?
3. How do various meta parameter searches contribute to get the best performance from these two models?

The research aims to compare the predictive performance metrics of logistic regression and support vector machine models for diabetes prediction based on patient data. It includes an in-depth analysis of the factors influencing the differences in performance metrics, with a specific focus on hyper parameters. Additionally, the study explores the impact of various meta-parameter searches on optimizing the overall performance of the models, aiming to identify the most effective parameters for diabetes prediction.

1.3 Aims and objectives

The primary aim of this study is to assess and compare the predictive performance metrics of logistic regression and support vector machine algorithms based on diabetic patient's data to decide if a patient is diabetic or not. This study aims to provide valuable insights into the strengths and limitations of these supervised machine learning approaches, contributing to the enhancement of diabetes prediction methodologies.

The main objectives of this study include:

- Obtain the model data, clean and analyze it and prepare it for model training.
- Train a standard logistic regression classification model on the labeled diabetes dataset.
- Train a competing SVM model on the same data.
- Explore model meta parameter and tune models to maximize predictive performance on accuracy, precision, recall, f1-score and confusion matrix.
- Evaluate the results of performance metrics after hyper parameter tuning to determine the best model for diabetes prediction.

1.4 Solution approach

The study follows a systematic approach encompassing model implementation, data preprocessing, splitting of data, hyper parameter tuning and performance analysis. Thorough exploration of various meta parameter search strategies contribute to achieving the aims. I will make sure to provide a clear documentation to make sure results are reliable and can be easily reproducible.

1.4.1 Dataset Description

This Diabetes Dataset Johndasilva (2018) has 2000 instances with 9 attributes - 'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'. The 'Outcome' attribute indicates positive or negative for diabetes.

1.4.2 Data Preprocessing

Data preprocessing is most important process. Mostly healthcare related data contains missing values that causes effectiveness of data. This process is essential for accurate result and successful predication using machine learning techniques.

Missing values removal

This process is meant to identify instances with zero value and eliminate all such instances. Through eliminating irrelevant instances we make feature subset and this process is called feature subset selection which reduces the dimensionality of data.

Splitting of data

After cleaning the data, the data is normalized in training and testing the models. After split, we train algorithm with training dataset and keep testing dataset aside. This testing dataset is used to test the trained model.

1.4.3 Model Implementation

After data is ready, we apply machine learning techniques. Implement logistic regression and support vector machine models to predict diabetes on the dataset.

Logistic Regression

Logistic Regression is one of the most common classification models. It is used for classification task where the goal is to predict the probability that an instance belongs to a given class or not. Create a standard logistic regression classification model, train the model with training dataset, and calculate the performance metrics. Similarly calculate the performance metrics for testing dataset. For this classifier, there are multiple hyper parameters such as regularization strength (C), solver, penalty.

Support Vector Machine

Support Vector Machine is a powerful machine learning algorithm used for linear or nonlinear classification, regression, and outlier detection tasks. This classifier aims to establish a hyperplane that can separate the classes by adjusting the distance between data points and the hyperplane.

1.4.4 Hyper Parameter Tuning and Performance Analysis

To determine the best performance model, we should consider the factors like hyper parameters, meta parameter search strategies for logistic regression and support vector machine to achieve more predictive performance metrics.

Grid Search

Grid Search is a method for hyper parameter optimization that involves specifying a list of values for each hyper parameter to optimize. Subsequently, the model is trained for each combination of these values, and the optimal values for the hyper parameters are selected based on the models' performance. For logistic regression classifier, there are multiple hyper parameters such as regularization strength (C), solver, penalty. For support vector machine classifier, there are multiple hyper parameters such as regularization parameter (C), kernel, gamma, degree.

1.5 Summary of contributions and achievements

This research focused on evaluating the performance of logistic regression and support vector machine models in predicting diabetes, examining the impact of hyperparameter tuning on these models. The primary objective was to compare the baseline performance of these models and then determine whether hyperparameter tuning could significantly improve their predictive capabilities.

Both models exhibited good performance on the training and testing dataset, while svm slightly outperformed lr on the test dataset, achieving 82% accuracy while lr achieved 79%. After applying hyperparameter tuning with gridsearchCV, svm achieved a training accuracy of 85% and testing accuracy of 80%, demonstrating the significant impact of tuning on model performance.

The results of this study have several implications:

- **Effectiveness of Hyperparameter Tuning:** The improved accuracy and other metrics demonstrate that hyperparameter tuning can lead to significant performance gains, making it a crucial step in machine learning model development.
- **Potential for SVM in Diabetes Prediction:** Given its superior performance, the svm model shows promise for further research and deployment in diabetes prediction applications.

Overall, this research has demonstrated that both models can effectively predict diabetes, and that hyperparameter tuning can significantly enhance their performance. Further research could explore additional tuning strategies, other machine learning models, or ensemble methods to improve prediction accuracy and reliability.

Chapter 2

Literature Review

The examination of relevant studies yields findings from various healthcare datasets, where researchers conducted analyses and predictions employing a range of methods and techniques. Numerous prediction models have been devised and applied by various researchers, utilizing different forms of data mining techniques, machine learning algorithms, or a combination of these methodologies.

2.1 Review of state-of-the-art

Mujumdar and Vaidehi (2019) aims to create a system using machine learning algorithm and deep learning techniques to provide accurate results and reduce human efforts. The diabetes dataset contains 800 instances with 10 attributes. This study implemented various machine learning algorithms include Support Vector Classifier, Random Forest Classifier, Decision Tree Classifier, Extra Tree Classifier, Ada Boost algorithm, Perceptron, Linear Discriminant Analysis algorithm, Logistic Regression, K-Nearest Neighbour, Gaussian Naïve Bayes, Bagging algorithm, Gradient Boost Classifier. The study incorporates the concept of pipelining and compares the diabetes dataset with the Pima dataset. Performance analysis includes metrics like classification accuracy, confusion matrix, f1-score, precision, and recall. The findings reveal that Logistic Regression achieves the highest accuracy of 96%, indicating an improvement in accuracy for the diabetes dataset compared to the Pima diabetes dataset. The study concludes that implementing a pipeline model enhances the accuracy of the classification performance, with the Ada Booster classifier identified as the best model, achieving an accuracy of 98.8%.

Soni and Varma (2020) aims to design and implement Diabetes prediction using machine learning methods and performance analysis of that methods for early prediction and to cure diabetes and save humans life. The diabetes dataset is gathered from UCI repository which is named as Pima Indian Diabetes dataset. The dataset have many attributes of 768 patients. The proposed methodology involves the utilization of various classification and ensemble learning methods, including SVM, Logistic Regression, KNN, Rndom Forest, Decision Tree, Gradient Boosting classifiers are used. The findings indicate a 77% accuracy achieved through an 80:20 split. The study concludes that Random Forest classifier exhibits highest accuracy when compared to other machine learning methods.

Swapna et al. (2018) aims to develop a methodology for classification of diabetic and normal Heart rate variability (HRV) signals using advanced deep learning architectures, specifically

employing Long Short-Term Memory (LSTM), Convolutional Neural Network (CNN), and combinations. The extracted features are then passed into a Support Vector Machine (SVM) for accurate classification. The study demonstrates performance improvements in CNN and CNN-LSTM architectures compared to earlier work, achieving a high accuracy of 95.7%.

2.2 Critique of the review

The literature review provides insights related to diabetes prediction using machine learning and deep learning techniques. Each study aims to contribute to the development of reliable and accurate models for diabetes diagnosis, with a focus on enhancing classification performance. Mujumdar and Vaidehi (2019) focus on creating an efficient system, achieving a 96% accuracy with Logistic Regression and identifying Ada Booster as the best model. Soni and Varma (2020) design a predictive system, achieving 77% accuracy with Random Forest being the most accurate. Swapna et al. (2018) explore advanced deep learning architectures, obtaining a high accuracy of 95.7% with CNN-LSTM and SVM. The key findings include the effectiveness of pipelining models, ensemble methods and the importance of accurate diabetes prediction for early intervention and patient care. However, the identified studies exhibit some limitations, such as a lack of detailed explanations regarding the selection of specific algorithms in the ensemble and a deficiency in exploring deep learning techniques that could enhance performance. Additionally, there is limited discussion on the generalizability of the proposed methodology to diverse datasets. Future research directions involve anomaly prediction and the utilization of larger datasets. Yudheksha et al. (2022) presented a machine learning-based approach for early-stage diabetes prediction using a dataset of patient attributes. They demonstrated the effectiveness of their model in predicting diabetes risk, achieving promising results. Additionally, Larabi-Marie-Sainte et al. (2019) conducted a comprehensive review of current techniques for diabetes prediction. Their review provided insights into various methods and strategies employed in diabetes prediction research, highlighting the importance of accurate prediction models in healthcare applications.

2.3 Summary

This literature review extensively examines the current research landscape in the realm of diabetes prediction, highlighting numerous technological advancements in this field. The literature review thoroughly investigates three prominent studies in diabetes prediction using machine learning and deep learning. Mujumdar and Vaidehi (2019) employ innovative pipelining techniques for accuracy enhancement but lack in-depth exploration of neural networks. Soni and Varma (2020) focus on diverse classifiers without extensive rationale, neglecting potential gains from deep learning. Swapna et al. (2018) Kamble and Patil (2016) excel in deep learning, emphasizing the need for a unified approach and comprehensive evaluation. The critiques call for collaboration, integration of machine learning strengths, and standardized methodologies for future research in diabetes prediction.

Chapter 3

Methodology

3.1 Dataset Description and Data Exploration

The diabetes data set is originated from kaggle (Johndasilva, 2018). The dataset contains 2000 patients and their corresponding 9 unique attributes. The nine attributes that are used for the prediction of diabetes are 'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'. The 'Outcome' attribute is taken as a dependent or target variable, and the remaining eight attributes are taken as independent feature variables.

Table 3.1: Dataset attributes and their data types

	Attribute	Description	Type
0	Pregnancies	Number of times pregnant	Numeric
1	Glucose	Plasma glucose concentration	Numeric
2	BloodPressure	Diastolic blood pressure	Numeric
3	SkinThickness	Triceps skinfold thickness	Numeric
4	Insulin	2-hour serum insulin	Numeric
5	BMI	Body mass index	Numeric
6	DiabetesPedigreeFunction	Diabetes pedigree function	Numeric
7	Age	Age (years)	Numeric
8	Outcome	Diabetes diagnose results (0: Negative, 1: Positive)	Nominal

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	2	138	62	35	0	33.6	0.127	47	1
1	0	84	82	31	125	38.2	0.233	23	0
2	0	145	0	0	0	44.2	0.630	31	1
3	0	135	68	42	250	42.3	0.365	24	1
4	1	139	62	41	480	40.7	0.536	21	0

Figure 3.1: Top 5 patients data

The diabetes attribute 'Outcome' is a categorical feature which consists of binary value where 0 means non-diabetes, and 1 implies diabetes. There are no null values for all attributes but there are zero values for few attributes which needs to be handled.

3.1.1 Data Distribution using Histograms

Fig 3.2 gives a better feel than the raw numbers and percentiles of the distributions of our numerical attributes. It shows how each feature and label is distributed along different ranges which further confirms the need for scaling. Next, wherever you see discrete bars, it basically means that each of these is actually a categorical variable. We will need to handle these categorical variables before applying Machine Learning. Our outcome labels have two classes, 0 for no disease and 1 for disease.




Figure 3.2: Data Distribution for each attribute

3.1.2 Correlation Matrix

A correlation matrix is a powerful tool for data analysis. It is a statistical technique used to evaluate the relationship between two variables in a dataset. It provides a correlation coefficient for each cell. The correlation coefficient value remains in the range between -1 and 1.

The correlation coefficient value -1 indicates notable negative linear correlation. The coefficient value 1 indicates notable positive linear correlation. The coefficient value 0 indicates no linear correlation.

A correlation matrix helps summarize data, identify patterns, and make decisions based on relationships between attributes. It helps us to gain insights for building better machine learning models by understanding which attributes are correlated.




Figure 3.3: Correlation Matrix

Here, we have to check the attributes correlated to target feature 'Outcome'. In Fig 3.3, 'Glucose' has highest correlation coefficient value - 0.46 and 'BMI' has 0.28, 'Age' has 0.24 and 'Pregnancies' has 0.22 compare to other attributes.

3.1.3 Bar plot for Outcome class

Fig 3.4 shows that the data is biased towards data points having outcome value as 0 which means that diabetes was not present actually. The number of non-diabetics is almost twice the number of diabetic patients.

Out of the 2000 instances, 1316 are associated with non-diabetic patients, while the remaining 684 pertain to diabetic patients. Therefore, it is essential to split the data efficiently for training and testing the model to ensure optimal results.

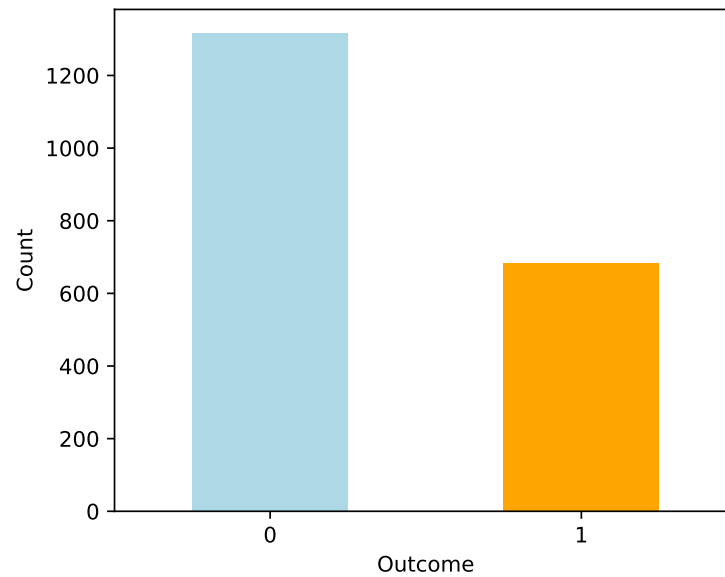


Figure 3.4: Distribution of Outcome (0s and 1s)

3.2 Data Preprocessing

Data preprocessing helps to transform data used to build a model which gives higher performance metrics. This process performs various functions like handling missing values, normalization and feature selection to improve the quality of data.

3.2.1 Missing Values Identification

There are no null values for all attributes. However, Fig 3.1 illustrates instances of zero values for attributes which are irrelevant and need to be handled. Table 3.2 shows the number of zero values for each attribute before imputation and after performing imputation.

Table 3.2: The number of zero missing values in dataset

Attributes	Count(Zeros)	Count(Zeros) after imputation
Pregnancies	301	0
Glucose	13	0
BloodPressure	90	0
SkinThickness	573	0
Insulin	956	0
BMI	28	0
DiabetesPedigreeFunction	0	0
Age	0	0

Here, we have observed numerous attributes with zero values, impacting the data quality. We replaced the zero values with the corresponding mean values using `simpleimputer`.

Fig 3.1 and Fig 3.5 shows the top five patients data. If you compare these figures, Fig 3.5 indicates that zero values in Fig 3.1 was replaced with mean for each attribute. Now, there are no missing values either null or zero values in each attribute.

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	2.000000	138.0	62.000000	35.000000	153.743295	33.6	0.127	47.0
1	4.359623	84.0	82.000000	31.000000	125.000000	38.2	0.233	23.0
2	4.359623	145.0	72.403665	29.341275	153.743295	44.2	0.630	31.0
3	4.359623	135.0	68.000000	42.000000	250.000000	42.3	0.365	24.0
4	1.000000	139.0	62.000000	41.000000	480.000000	40.7	0.536	21.0

Figure 3.5: Top 5 patients data after handling missing values

3.2.2 Feature Selection based on Correlation Coefficient

After handling missing values, the correlation coefficient is calculated in this method which correlates with the output and input attributes. This was discussed under section 3.1.2. We have created a Table 3.3 to show correlation coefficient values of each attribute towards target attribute.

Table 3.3: The correlation coefficient values

Attributes	Correlation Coefficient	Correlation Coefficient after imputation
Pregnancies	0.224437	0.249883
Glucose	0.458421	0.488020
BloodPressure	0.075958	0.174481
SkinThickness	0.076040	0.205527
Insulin	0.120924	0.207696
BMI	0.276726	0.282182
DiabetesPedigreeFunction	0.155459	0.155459
Age	0.236509	0.236509

We used 0.2 as a cut-off for relevant attributes. Hence 'BloodPressure' and 'DiabetesPedigreeFunction' features are removed. 'Pregnancies', 'Glucose', 'SkinThickness', 'Insulin', 'BMI', and 'Age' are our most relevant six input attributes.

3.2.3 Data Normalization

Normalization refers to the process of scaling and transforming numeric features to a standard scale or distribution. Feature scaling is a technique to standardize or normalize the range of independent features or variables of a dataset. The goal of feature scaling is to ensure that all features contribute equally to the learning process, preventing certain features from dominating others based on their scale. In Table 3.3, we can see that 'Glucose' and 'Outcome' have a 0.49 correlation coefficient. Hence these are highly correlated. After completing data preprocessing, we have total 2000 instances.

3.3 Dataset Split into Train and Test Data

After data cleaning and preprocessing, the dataset becomes ready to train and test. We are using test split method to split data randomly into the training and testing set. Here, I am partitioning the data into a training set comprising 70% and a test set comprising 30%.

3.4 Model Implementation

This is the important phase which includes model building for diabetes prediction.

3.4.1 Logistic Regression

It is used for classification task where the goal is to predict the probability that an instance belongs to a given class or not.

Algorithm 1 Diabetes Prediction using Logistic Regression

Input: Input features X and labels y for training data

Output: Generate performance metrics like accuracy, confusion matrix, precision, recall, f1-score

- 1: Create a standard logistic regression model with default hyperparameters such as regularization strength (C), solver, penalty
 - 2: Fit the model with training data
 - 3: Calculate accuracy, confusion matrix, precision, recall and f1-score for the trained model
 - 4: Predict outcomes for testing data using the trained model
 - 5: Evaluate the model performance on testing data
 - 6: Calculate accuracy, confusion matrix, precision, recall and f1-score
-

3.4.2 Support Vector Machine

It is used for linear or nonlinear classification, regression, and outlier detection tasks. This classifier aims to establish a hyperplane that can separate the classes by adjusting the distance between data points and the hyperplane.

Algorithm 2 Diabetes Prediction using Support Vector Machine

Input: Input features X and labels y for training data

Output: Generate performance metrics like accuracy, confusion matrix, precision, recall, f1-score

- 1: Create a standard svm model with default hyperparameters such as regularization parameter (C), kernel, gamma, degree
 - 2: Fit the model with training data
 - 3: Calculate accuracy, confusion matrix, precision, recall and f1-score for the trained model
 - 4: Predict outcomes for testing data using the trained model
 - 5: Evaluate the model performance on testing data
 - 6: Calculate accuracy, confusion matrix, precision, recall and f1-score
-

3.4.3 Hyperparameter Tuning with GridSearchCV

Hyperparameter Tuning is a process to select optimal values for a machine learning models hyperparameters. The goal of hyperparameter tuning is to find the values that leads to the best performance for a given problem. We should consider the factors like hyper parameters, meta parameter search strategies to achieve more predictive performance metrics.

Grid Search

Grid Search is a method for hyper parameter optimization that involves specifying a list of values for each hyper parameter to optimize. Subsequently, the model is trained for each combination of these values, and the optimal values for the hyper parameters are selected based on the models' performance.

Hyperparameter Tuning for Logistic Regression

Logistic Regression is one of the most common classification algorithms. It is used for classification task where the goal is to predict the probability that an instance belongs to a given class or not. There are multiple hyper parameters such as regularization strength (C), solver, penalty. We have solvers namely, 'lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga'. We have penalties namely, 'l1', 'l2', 'elasticnet', 'none'.

Algorithm 3 Diabetes Prediction using hyperparameter tuning technique for Logistic Regression

Input: Input features X and labels y for training data

Output: Generate best parameters and calculate performance metrics like accuracy, confusion matrix, precision, recall, f1-score using best parameters

- 1: Create a standard logistic regression model with no hyperparameters such as regularization strength (C), solver, penalty
 - 2: Create a paramgrid dictionary with all hyperparameters related to logistic regression model you would like to pass to gridsearch for tuning
 - 3: Create a gridsearchcv model by passing parameters such as model estimator, paramgrid, cross validation, scoring methods, refit, verbose
 - 4: Generate the best estimators, best parameters and best scores for training set
 - 5: Calculate accuracy, confusion matrix, precision, recall and f1-score for trained model
 - 6: Predict outcomes for testing data using the trained model
 - 7: Evaluate the model performance on testing data
 - 8: Calculate accuracy, confusion matrix, precision, recall and f1-score
-

Hyperparameter Tuning for Support Vector Machine

This classifier aims at forming a hyper plane that can separate the classes as much as possible by adjusting the distance between the data points and the hyper plane. There are several kernels based on which the hyper plane is decided. There are multiple hyper parameters such as regularization parameter (C), kernel, gamma, degree. We have kernels namely, 'linear', 'poly', 'rbf', and 'sigmoid'.

Algorithm 4 Diabetes Prediction using hyperparameter tuning technique for Support Vector Machine

Input: Input features X and labels y for training data

Output: Generate performance metrics like accuracy, confusion matrix, precision, recall, f1-score

- 1: Create a standard svm model with no hyperparameters such as regularization parameter (C), kernel, gamma, degree
 - 2: Create a paramgrid dictionary with all hyperparameters related to svm model you would like to pass to gridsearch for tuning
 - 3: Create a gridsearchcv model by passing parameters such as model estimator, paramgrid, cross validation, scoring methods, refit, verbose
 - 4: Generate the best estimators, best parameters and best scores for training set
 - 5: Calculate accuracy, confusion matrix, precision, recall and f1-score for trained model
 - 6: Predict outcomes for testing data using the trained model
 - 7: Evaluate the model performance on testing data
 - 8: Calculate accuracy, confusion matrix, precision, recall and f1-score
-

3.5 Evaluation Metrics

Evaluation would be the final step of prediction model. Here, we evaluate the prediction results using various evaluation metrics like classification accuracy, confusion matrix, precision, recall, f1-score.

3.5.1 Classification Accuracy

It is the ratio of number of correct predictions to the total number of input samples.

$$Accuracy = \frac{\text{Number of correct predictions}}{\text{Total number of predictions made}} \quad (3.1)$$

3.5.2 Confusion Matrix

It provides us a matrix output that describes the performance of the model.

$$\text{Confusion Matrix} = \begin{pmatrix} TP & FP \\ FN & TN \end{pmatrix} \quad \begin{array}{l} \text{Where : TP = True Positives} \\ \text{FP = False Positives} \\ \text{FN = False Negatives} \\ \text{TN = True Negatives} \end{array} \quad (3.2)$$

3.5.3 Precision

It is the number of correct positive results divided by number of positive results predicted by the classifier.

$$Precision = \frac{TP}{(TP + FP)} \quad (3.3)$$

3.5.4 Recall

It is the number of correct positive results divided by number of all relevant samples.

$$Recall = \frac{TP}{(TP + FN)} \quad (3.4)$$

3.5.5 F1-score

It is used to measure a test's accuracy. F1-score is the harmonic mean between precision and recall. The range for f1-score is [0, 1]. It tells you how precise your classifier is as well as how robust it is.

$$F1 = 2 \times \frac{1}{\left(\frac{1}{precision}\right) + \left(\frac{1}{recall}\right)} \quad (3.5)$$

3.6 Summary

This methodology provides a structured approach to exploring and showcasing the significance of hyperparameter tuning in enhancing the predictive capabilities of logistic regression and support vector machine models for diabetes prediction. This structured approach encompasses aspects such as dataset characteristics, data exploration, data visualization, preprocessing, normalization, and feature selection to ensure data quality and model design. It delves into data splitting for training and testing, the training of classification models using the training dataset, and the crucial step of hyperparameter tuning to enhance performance metrics by identifying optimal parameters.

Chapter 4

Results

4.1 Results for ML methods - LR and SVM

The evaluation of machine learning algorithms is crucial for assessing their effectiveness in predicting diabetes. We use various performance metrics such as accuracy, precision, recall and f1-score to compare the classification methods. These metrics are calculated using equations 3.1, 3.3, 3.4, 3.5, derived from the confusion matrix.

Table 4.1 shows performance measures of logistic regression and svm classifiers for training and testing datasets.

Table 4.1: Performance metrics

Classifier	DataSet Type	Precision	Recall	F1-score	Accuracy (%)
LR	Training	0.70	0.56	0.62	77%
LR	Testing	0.75	0.58	0.65	79%
SVM	Training	0.76	0.71	0.74	83%
SVM	Testing	0.77	0.66	0.71	82%

Here, lr model performed adequately on the training and testing datasets while svm model achieved higher accuracy, with 83% on the training set and 82% on the testing set. This higher accuracy, along with better precision, recall, and F1-score, suggests that svm might be a more suitable model for diabetes prediction. The confusion matrix for the svm model on the test dataset is provided in Table 4.2.

Table 4.2: Confusion Matrix for SVM (Test Data)

	Diabetic	Non-diabetic
Diabetic	236	27
Non-diabetic	46	91

Despite these promising results, determining the best model for predicting diabetes requires further investigation and analysis.

4.2 Hyperparameter Tuning Results for LR and SVM Using GridSearchCV

After applying gridsearchCV for hyperparameter tuning, both logistic regression and svm models demonstrated significantly improved performance compared to their standard configurations which was shown in Table 4.1.

Table 4.3 shows performance measures of logistic regression and svm classifier for training and testing datasets after tuning.

Table 4.3: Performance metrics after hyperparameter tuning

Classifier	DataSet Type	Precision	Recall	F1-score	Accuracy (%)
LR	Training	0.72	0.54	0.62	77%
LR	Testing	0.80	0.57	0.67	81%
SVM	Training	0.79	0.76	0.77	85%
SVM	Testing	0.78	0.66	0.72	82%

Hyperparameter tuning significantly improved both models' performance. Here, svm model achieved the highest accuracy 85% on the training dataset and 82% on the testing dataset. LR model also saw improvements, with 81% testing accuracy. While both models showed good precision, the recall could still be enhanced. Fine-tuning hyperparameters led to improved overall accuracy and balanced performance. However, further refinement may be needed to increase recall and reduce misclassifications. The confusion matrix for the svm model on the test dataset is provided in Table 4.4

Table 4.4: Confusion Matrix for SVM (Test Data)

	Diabetic	Non-diabetic
Diabetic	238	25
Non-diabetic	46	91

The optimized models achieved higher accuracy scores, highlighting the effectiveness of hyperparameter tuning in enhancing predictive capabilities.

4.3 Summary

The evaluation of machine learning algorithms lr and svm for predicting diabetes demonstrated promising results. Although both models had comparable performance on the training dataset, svm outperformed lr on the testing dataset, achieving 82% accuracy. Hyperparameter tuning with gridsearchCV further improved the models' performance, with the optimized svm achieving an accuracy of 82% on the testing set, highlighting the effectiveness of this technique. These results suggest that svm might be more effective for diabetes prediction, but additional tuning and exploration could further enhance the performance of both models. These findings underscore the effectiveness of hyperparameter tuning in improving predictive capabilities.

Chapter 5

Discussion and Analysis

5.1 Discussion on performance metrics

The evaluation of machine learning algorithms lr and svm for predicting diabetes provides valuable insights into their effectiveness in healthcare applications. The results indicate that both models achieved commendable accuracy scores on the testing dataset, with svm showing a slight advantage over lr. We can discuss the implications of these findings and provides a deeper analysis of their significance.

5.2 Significance of the findings

The significance of the findings lies in their contribution to advancing the field of predictive analytics and healthcare informatics. By evaluating lr and svm models for predicting diabetes, this research provides valuable insights into the efficacy of machine learning algorithms in healthcare decision-making. Accurate prediction of diabetes can aid healthcare professionals in early detection and intervention, thereby improving patient management and reducing the risk of complications.

Firstly, the findings underscore the importance of leveraging advanced analytics techniques to extract actionable insights from healthcare data. By demonstrating the effectiveness of lr and svm models in predicting diabetes onset, this research validates the utility of machine learning approaches in healthcare research and practice.

Secondly, the findings contribute to the methodological advancement of predictive modeling in healthcare. By evaluating the performance of lr and svm models using various performance metrics such as accuracy, precision, recall, and f1-score, this research provides a comprehensive assessment of model performance. Moreover, by conducting hyperparameter tuning using gridsearchCV, the study demonstrates the importance of optimizing model parameters to enhance predictive accuracy and generalizability.

5.3 Limitations

The presented results provide valuable insights into the performance of logistic regression and support vector machine models for predicting diabetes. However, there are several key limitations and potential implications or improvements that should be considered:

- **Limited Model Comparison:** The analysis focuses only on lr and svm models, including a broader range of machine learning algorithms such as random forests, gradient boosting machines, or neural networks could reveal superior approaches and enhance understanding of model performance.
- **Evaluation Metrics:** Standard metrics like accuracy, precision, recall, and F1-score may not fully capture model performance, especially in scenarios with class imbalance or varied misclassification costs. Incorporating metrics like AUC-ROC or cost-sensitive evaluation measures can provide deeper insights and address these limitations.
- **Hyperparameter Tuning Sensitivity:** While hyperparameter tuning improves model performance, the results may be sensitive to the choice of hyperparameters and the specific hyperparameter search space. Conducting sensitivity analyses or exploring alternative hyperparameter optimization techniques could provide insights into the robustness of the tuned models.
- **Feature Selection:** Correlation matrix-based feature selection may overlook non-linear relationships with the target variable and assumes linear relationships between features. Employing advanced techniques like recursive feature elimination with cross-validation (RFECV) or tree-based methods can better capture non-linear relationships and improve feature selection accuracy.

5.4 Summary

The discussion and analysis chapter provides a comprehensive overview of the evaluation of machine learning algorithms lr and svm for predicting diabetes. It highlights the significance of the findings in advancing predictive analytics in healthcare, emphasizing the importance of leveraging advanced analytics techniques and optimizing model parameters for accurate predictions. However, it acknowledges several limitations, including the need for broader model comparisons, consideration of additional evaluation metrics, sensitivity to hyperparameter tuning, and refinement of feature selection techniques. There is room for improvement and further exploration of alternative methodologies to enhance predictive accuracy and robustness in healthcare applications.

Chapter 6

Conclusions and Future Work

6.1 Conclusions

Early detection of diabetes is one of the significant challenges in the health care industry. Our research delved into the realm of diabetes prediction using logistic regression and support vector machines models, aiming to shed light on their comparative performance metrics and underlying factors.

In our research, we designed a system which can predict diabetes with high accuracy. We pre-processed the data by addressing missing values, zero-valued features, and employing imputation techniques across all features. Using the feature reduction method, we dropped two features. We used six input features - 'Pregnancies', 'Glucose', 'SkinThickness', 'Insulin', 'BMI' and 'Age' and one output feature 'Outcome' in the dataset (Johndasilva, 2018). We explored the efficacy of two machine learning algorithms, lr and svm to predict diabetes and evaluated the performance on various measures like accuracy, precision, recall, and f1-score. These models provided an accuracy greater than 75%.

Furthermore, our exploration of hyperparameter optimization strategies underscores the importance of fine-tuning model parameters to achieve optimal predictive performance. To enhance model performance, we employed hyperparameter tuning using gridsearchCV strategy, optimizing parameters. These models provided good accuracy before tuning and svm outperformed lr for both train/test split method. After tuning, the optimized models exhibited notable improvements in accuracy, with svm achieving an impressive accuracy of approximately 82% on the test dataset. while lr slightly trailing behind, still demonstrated commendable performance, boasting an accuracy close to 81%.

6.2 Future work

The results from Chapter 4 indicate that both logistic regression and support vector machine classifiers achieved reasonable performance in predicting diabetes, with svm attaining slightly higher accuracy. However, there are several areas for future research that can further enhance the predictive capabilities of these models and address their limitations.

- A critical aspect for future work is the expansion of the dataset. The current study dataset may not fully represent the diverse population affected by diabetes. Increasing the dataset's

size and diversity would improve the model's generalizability and robustness. It would also help reduce biases and ensure that the models work across different demographics and geographic locations.

- Future studies can delve into additional feature engineering techniques. By incorporating domain-specific knowledge from medical professionals and using automated feature selection methods like recursive feature elimination, new features may be uncovered that contribute significantly to prediction accuracy. This approach could also lead to more refined models with reduced risks of overfitting.
- Future work could explore more complex architectures such as deep neural networks (DNNs), convolutional neural networks (CNNs), and recurrent neural networks (RNNs). These models might capture intricate patterns in the data that traditional machine learning models might miss. Deep learning methods could improve prediction accuracy and add robustness to the models.
- Another crucial aspect for future work is managing class imbalances. Class imbalance in the dataset can lead to biased models and high false negatives. Techniques like oversampling, undersampling, and Synthetic Minority Over-sampling Technique (SMOTE) could be used to balance the classes, resulting in improved recall and better model performance.
- The current study utilized `gridsearchCV` for hyperparameter tuning, other methods such as `randomizedsearchcv` and bayesian optimization offer more flexible and efficient tuning processes. Future work could investigate these methods to optimize model performance further, potentially achieving higher accuracy and reliability.
- Integrating these models into clinical practice is a valuable future direction. Building a user-friendly interface for healthcare professionals and conducting clinical trials would be essential to validate the models' effectiveness in real-world. This step is crucial for demonstrating that these models can assist in early diabetes diagnosis and improve patient outcomes.

Chapter 7

Reflection

Throughout this project, I gained significant skills in technical problem-solving. The diabetes prediction work involved diving into machine learning, experimenting with logistic regression and support vector machine algorithms, and mastering data processing. Although I had some prior knowledge, the in-depth exploration of model evaluation, optimization, and tuning proved to be an invaluable learning experience.

A crucial skill I developed during this project was the ability to systematically tackle complex problems. My initial plan was straightforward: gather data, preprocess it, and apply machine learning to predict diabetes. However, I faced unexpected challenges, including selecting performance metrics, choosing features, and tuning hyperparameters. These hurdles led me to adopt an iterative approach, emphasizing continuous improvement and validating each step to ensure the best outcomes.

This iterative process enhanced my skills in experimenting with `lr` and `svm` algorithms, comparing their performance, and fine-tuning them for better results. Learning to use `gridsearchCV` for hyperparameter tuning was a key takeaway—it showed me how even small parameter changes could greatly impact model accuracy and generalization.

Despite these successes, I faced challenges I couldn't fully resolve. Data imbalance and balancing precision with recall while keeping high accuracy were particularly tricky. I learned that perfect accuracy is often unattainable and that choosing the right model involves trade-offs. This experience deepened my understanding of machine learning and highlighted the need to evaluate models from multiple perspectives.

If I could redo this project, I would start with a thorough examination of the data, addressing imbalance earlier, and consider a wider variety of algorithms. I would also implement more rigorous cross-validation techniques from the outset to ensure my results were reliable and not influenced by random data fluctuations.

One deviation from my initial plan was the extended focus on hyperparameter tuning. Initially, I underestimated the impact it could have on model performance. As I progressed, I realized that tuning was essential for achieving optimal results. This shift in focus taught me the importance of flexibility and adaptability in research.

Overall, this project was a profound learning experience that enhanced my technical skills and problem-solving abilities. It also provided insights into the complexities of machine learning and the need for careful planning and adaptation. The lessons learned will undoubtedly influence my future work and approach to problem-solving in a broader context.

References

Johndasilva (2018), 'Diabetes dataset'. (accessed January 25, 2024).

URL: <https://www.kaggle.com/datasets/johndasilva/diabetes>

Kamble, M. T. P. and Patil, S. (2016), 'Diabetes detection using deep learning approach', *International Journal for Innovative Research in Science & Technology* **2**(12), 342–349.

Larabi-Marie-Sainte, S., Aburahmah, L., Almohaini, R. and Saba, T. (2019), 'Current techniques for diabetes prediction: review and case study', *Applied Sciences* **9**(21), 4604.

Mujumdar, A. and Vaidehi, V. (2019), 'Diabetes prediction using machine learning algorithms', *Procedia Computer Science* **165**, 292–299.

Soni, M. and Varma, S. (2020), 'Diabetes prediction using machine learning techniques', *International Journal of Engineering Research & Technology (Ijert)* Volume **9**.

Swapna, G., Vinayakumar, R. and Soman, K. (2018), 'Diabetes detection using deep learning algorithms', *ICT express* **4**(4), 243–246.

Yudheksha, G., Murugadoss, V., Reddy, P. S., Harshavardan, T. and Sriramulu, S. (2022), A machine learning based approach to early stage diabetes prediction, *in* '2022 International Conference on Automation, Computing and Renewable Systems (ICACRS)', IEEE, pp. 1275–1280.

Appendix A

An Appendix Chapter (Optional)

Some lengthy tables, codes, raw data, length proofs, etc. which are **very important but not essential part** of the project report goes into an Appendix. An appendix is something a reader would consult if he/she needs extra information and a more comprehensive understating of the report. Also, note that you should use one appendix for one idea.

An appendix is optional. If you feel you do not need to include an appendix in your report, avoid including it. Sometime including irrelevant and unnecessary materials in the Appendices may unreasonably increase the total number of pages in your report and distract the reader.

Appendix B

An Appendix Chapter (Optional)

...